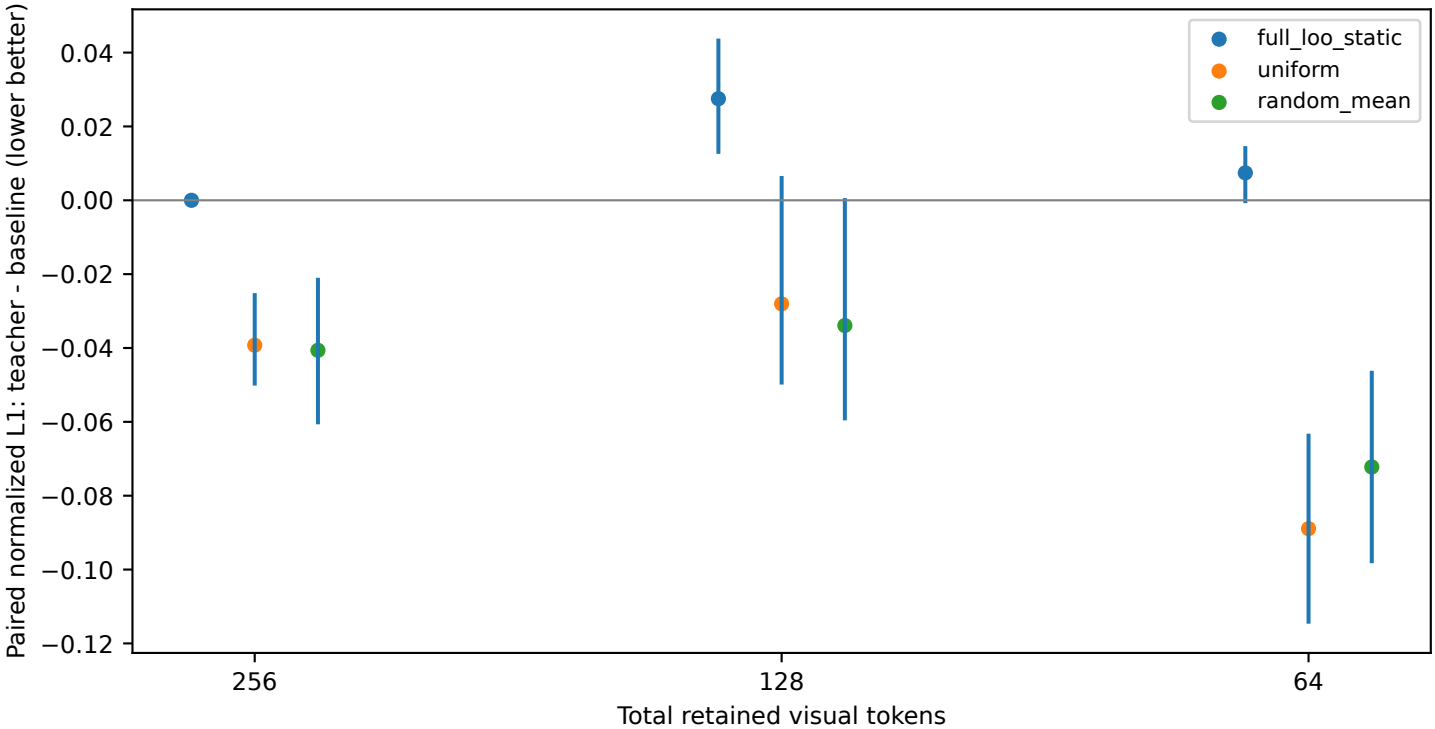
Teacher path: 12 observations / 4 episode clusters
95% episode bootstrap; exploratory, no automatic go
full_loo_static
uniform
random_mean
Paired normalized L1: teacher - baseline (lower better)
0.04
0.02
0.00
−0.02
−0.04
−0.06
−0.08
−0.10
−0.12
256
128
64
Total retained visual tokens